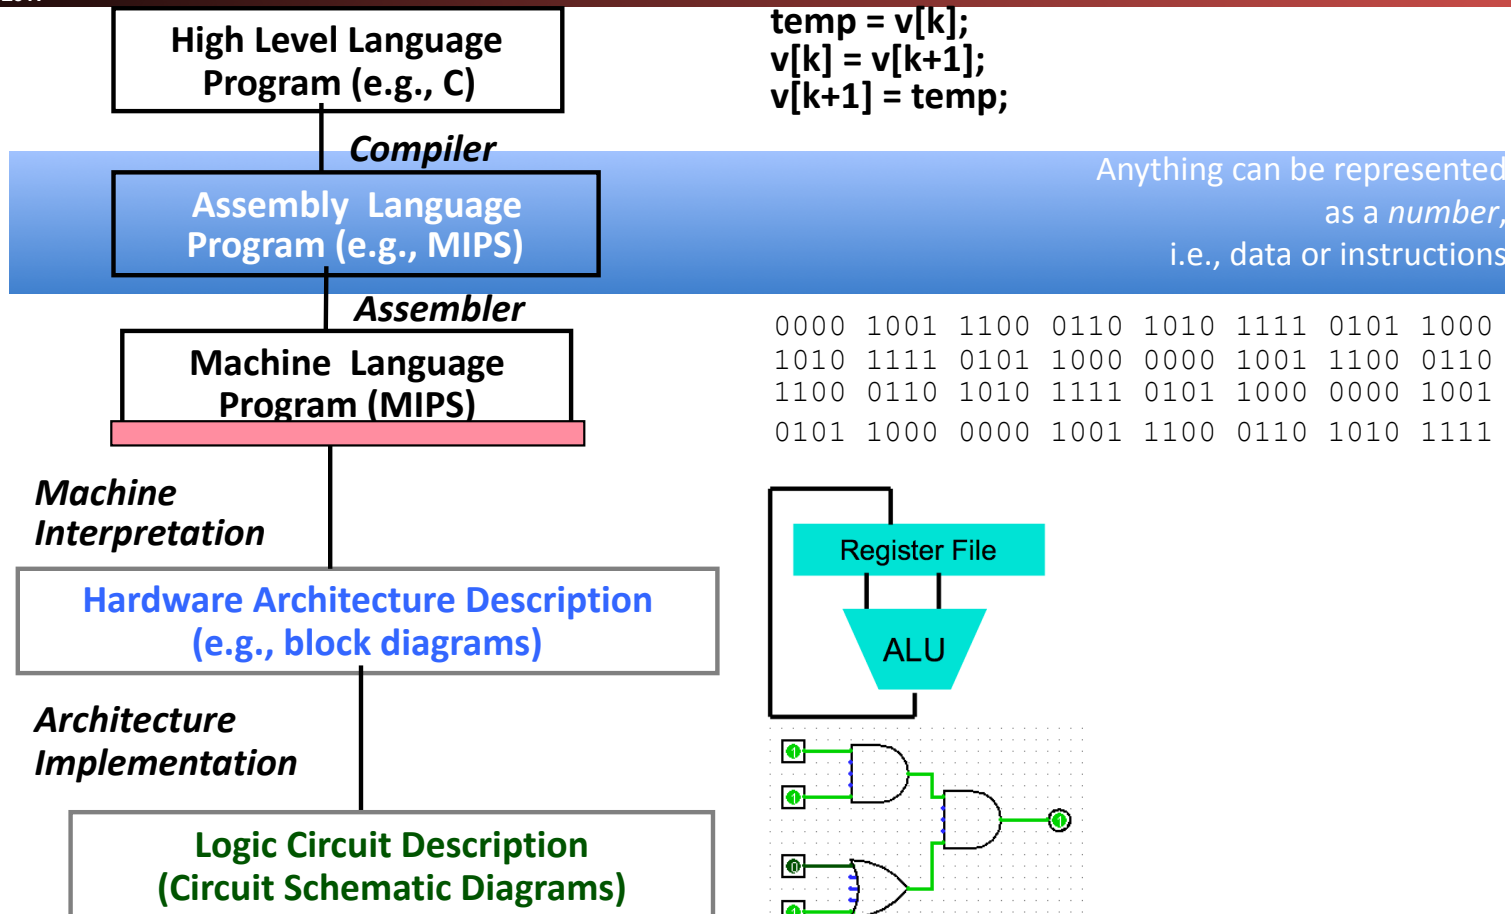


CS 61C:

Great Ideas in Computer Architecture

Intro to Assembly Language, MIPS Intro

Levels of Representation/Interpretation




Assembly Language

- Basic job of a CPU: execute lots of *instructions*.
- Instructions are the primitive operations that the CPU may execute
 - The stored program in "stored program computer"
- Different CPUs implement different sets of instructions. The set of instructions a particular CPU implements is an Instruction Set Architecture (ISA).
 - Examples: ARM, Intel x86, MIPS, RISC-V, IBM/Motorola PowerPC (old Mac), Motorola 68k (really old Mac), Intel IA64 (aka Itanic), ...

Instruction Set Architectures

- Early trend was to add more and more instructions to new CPUs to do elaborate operations
 - VAX architecture had an instruction to multiply polynomials!
- RISC philosophy (Cocke IBM, Patterson, Hennessy, 1980s)
Reduced Instruction Set Computing
 - Keep the instruction set small and simple, makes it easier to build fast hardware
 - Let software do complicated operations by composing simpler ones

Berkeley Acknowledge for the first RISC computer



From RISC-I to RISC-V

The RISC-V Instruction Set Architecture

RISC-V (pronounced "risk-five") is a new instruction set architecture (ISA) that was originally designed to support computer architecture research and education, which we now hope will become a standard open architecture for industry implementations. RISC-V was originally developed in the [Computer Science Division](#) of the EECS Department at the [University of California, Berkeley](#).

EECS151/251A, CS152, CS250, CS252 use RISC-V

But at the same time... All RISCs are mostly the same except for one or two silly design decisions

- MIPS:
 - "The Branch Delay Slot": The instruction immediately after an if is always executed
- SPARC:
 - "Register Windows"
- ARM:
 - 32b is ugly as sin...
 - But 64b is pretty much generic RISC with condition codes for instructions: Allow you to say "only do this instruction if a previous value was TRUE/FALSE"

MIPS Architecture

- MIPS – semiconductor company that built one of the first commercial RISC architectures
- We will study the MIPS architecture in some detail in this class
 - When an upper division course uses RISC-V instead, the ISA is very similar)
- Why MIPS instead of Intel x86?
 - MIPS is simple, elegant. Don't want to get bogged down in gritty details.
 - MIPS widely used in embedded apps, x86 little used in embedded, and more embedded computers than PCs
 - And you are more likely to need C and/or assembly in dealing with embedded systems
- Why MIPS instead of ARM?
 - Accident of history: Patterson & Hennessey is in MIPS
 - MIPS also designed more for performance than ARM, ARM is instead designed for small code size

Assembly Variables: Registers

- Unlike HLL like C or Java, assembly cannot use variables
 - Why not? Keep Hardware Simple
- Assembly Operands are ***registers***
 - Limited number of special locations built directly into the hardware
 - Operations can only be performed on these!
- Benefit: Since registers are directly in hardware, they are very fast
(faster than 1 ns - light travels 30cm in 1 ns!!!)

Number of MIPS Registers

- Drawback: Since registers are in hardware, there are a predetermined number of them
 - Solution: MIPS code must be very carefully put together to efficiently use registers
- 32 registers in MIPS
 - Why 32? Smaller is faster, but too small is bad. Goldilocks problem.
 - Plus can use 5 bits to specify the register
 - X86 has many fewer registers
 - Under the hood, a modern x86 actually translates to an internal RISC!
- Each MIPS register is 32 bits wide
 - Groups of 32 bits called a word in MIPS

Names of MIPS Registers

- Registers are numbered from 0 to 31
- Each register can be referred to by number or name
- Number references:
 - \$0, \$1, \$2, ... \$30, \$31
- For now:
 - \$16 - \$23 → \$s0 - \$s7 (correspond to C variables)
 - \$8 - \$15 → \$t0 - \$t7 (correspond to temporary variables)
 - Later will explain other 16 register names
- In general, use register names to make your code more readable

C, Java variables vs. registers

- In C (and most High Level Languages) variables declared first and given a type
 - Example: `int fahr, celsius; char a, b, c, d, e;`
- Each variable can ONLY represent a value of the type it was declared
 - Can't treat an `int` like a `char`, and converting (casting) a `char` to an `int` changes type.
- In Assembly, registers have no type:
 - The operation determines how to treat the register's contents

Addition and Subtraction of Integers

- Addition in Assembly
 - MIPS: `add $s0,$s1,$s2`
C: `a = b + c`
 - where C variables \Leftrightarrow MIPS registers are:
`a \Leftrightarrow $s0, b \Leftrightarrow $s1, c \Leftrightarrow $s2`
- Subtraction in Assembly
 - MIPS: `sub $s0,$s1,$s2`
C: `a = b - c`
 - where C variables \Leftrightarrow MIPS registers are:
`a \Leftrightarrow $s0, b \Leftrightarrow $s1, c \Leftrightarrow $s2`

Addition and Subtraction of Integers Example 1

- How to do the following C statement?
a = b + c + d - e;
- Break into multiple instructions
 - **add \$t0, \$s1, \$s2 # temp = b + c**
add \$t0, \$t0, \$s3 # temp = temp + d
sub \$s0, \$t0, \$s4 # a = temp - e
- A single line of C may break up into several lines of MIPS.
- Notice the use of temporary registers – don't want to modify the variable registers \$s
- Everything after the hash mark on each line is ignored (comments)

Immediates

- Immediates are numerical constants
 - They appear often in code, so there are special instructions for them
- Add Immediate:
 - `addi $s0,$s1,-10` (in MIPS)
 - `f = g - 10` (in C)
 - where MIPS registers `$s0`, `$s1` are associated with C variables `f`, `g`
- Syntax similar to add instruction, except that last argument is a number instead of a register

Overflow in Arithmetic

- Reminder: Overflow occurs when there is a “mistake” in arithmetic due to the limited precision in computers.

- Example (4-bit unsigned numbers):

15	1111
+ 3	+ 0011
18	10010

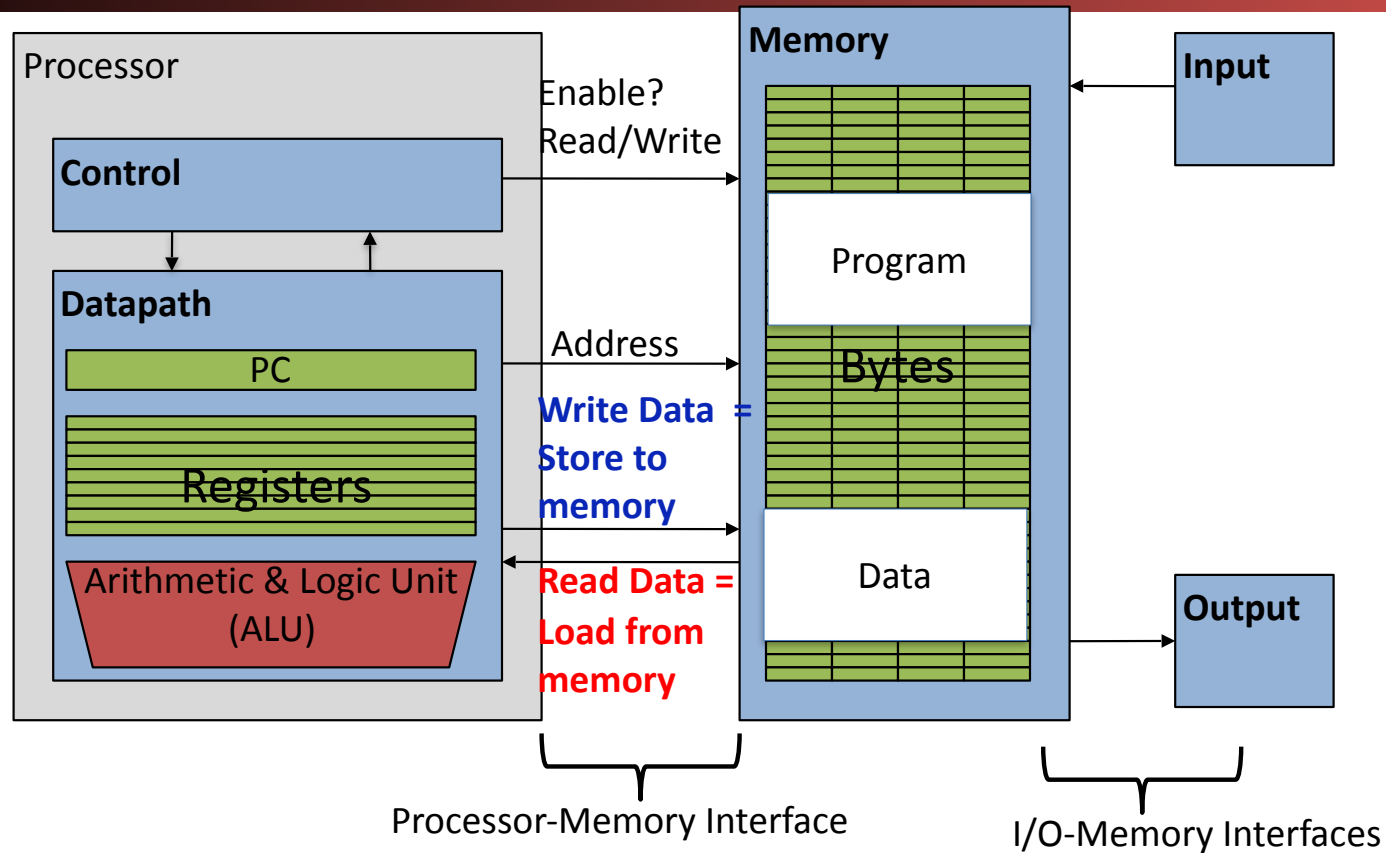
- But we don't have room for 5-bit solution, so the solution would be 0010, which is +2, and “wrong”.

Overflow handling in MIPS

- Some languages detect overflow (Ada), some don't (C, Java)
- MIPS solution is 2 kinds of instructions
 - The **only** difference is what happens on overflow!
- "signed" cause overflow to be detected
 - add (**add**)
 - add immediate (**addi**)
 - subtract (**sub**)
- "unsigned" do not cause overflow detection
 - add unsigned (**addu**)
 - add immediate unsigned (**addiu**)
 - subtract unsigned (**subu**)
- Compiler selects appropriate arithmetic
 - MIPS C compilers produce **addu**, **addiu**, **subu**

Data Transfer:


Load from and **Store to** memory



Memory Addresses are in Bytes

- Lots of data is smaller than 32 bits, but rarely smaller than 8 bits – works fine if everything is a multiple of 8 bits
- 8 bit chunk is called a byte
 - (1 MIPS word = 4 bytes, a "32b system")
- Memory addresses are really in bytes, not words
- Word addresses are 4 bytes apart
 - Word address is same as address of leftmost byte – most significant byte (i.e. Big-endian convention)

Most significant byte in a word



...
12	13	14	15
8	9	10	11
4	5	6	7
0	1	2	3

Transfer from Memory to Register


- C code:
 - `int A[100];`
`g = h + A[3];`
- Using Load Word (lw) in MIPS:
 - `lw $t0,12($s3) # Temp reg $t0 gets A[3]`
`add $s1,$s2,$t0 # g = h + A[3]`
- Note: `$s3 == A` – base register (pointer)
- `12 == offset in bytes`
 - Offset must be a constant known at assembly time

Transfer from Register to Memory

- C code
 - `int A[100];`
`A[10] = h + A[3];`
- Using Store Word (sw) in MIPS:
 - `lw $t0, 12($s3) # Temp reg $t0 gets A[3]`
 - `add $t0, $s2, $t0 # Temp reg $t0 gets h + A[3]`
 - `sw $t0, 40($s3) # A[10] = h + A[3]`
- `$s3+12` and `$s3+40` must be multiples of 4: Word alignment!

Loading and Storing bytes

- In addition to word data transfers (**lw**, **sw**), MIPS has byte data transfers:
 - load byte: **lb**
 - store byte: **sb**
- Same format as **lw**, **sw**
 - But offsets can be individual addresses, they don't have to be word aligned
- E.g., **lb \$s0, 3(\$s1)**
 - contents of memory location with address = sum of “3” + contents of register \$s1 is copied to the low byte position of register \$s0.



How Much Progress on the Project?



A: What's a project?

B: Seen Piazza Post

C: Gotten Code, started

D: Done with rebugging (aka "Programming")

E: Done with debugging

Clickers/Peer Instruction

We want to translate $*x = *y + 1$ into MIPS
(x, y pointers stored in: $\$s0$ $\$s1$)

A: `addi $s0, $s1, 1`

B: `lw $s0, 1($s1)`
`sw $s1, 0($s0)`

C: `lw $t0, 0($s1)`
`addi $t0, $t0, 1`
`sw $t0, 0($s0)`

D: `sw $t0, 0($s1)`
`addi $t0, $t0, 1`
`lw $t0, 0($s0)`

E: `lw $s0, 1($t0)`
`sw $s1, 0($t0)`

Speed of Registers vs. Memory

- Given that
 - Registers: 32 words (128 Bytes)
 - Memory: Billions of bytes (2 GB to 8 GB on laptop)
- and the RISC principle is...
 - Smaller is faster
- How much faster are registers than memory??
- About 100-500 times faster!
 - in terms of latency of one access
 - We will later learn about caches to reduce this latency in practice

MIPS Logical Instructions

- Useful to operate on fields of bits within a word
 - e.g., characters within a word (8 bits)
- Operations to pack /unpack bits into words
- Called *logical operations*

Logical operations	C operators	Java operators	MIPS instructions
Bit-by-bit AND	&	&	and
Bit-by-bit OR			or
Bit-by-bit NOT	~	~	not
Shift left	<<	<<	sll
Shift right	>>	>>>	srl

Logic Shifting

- Shift Left:
`sll $s1,$s2,2 #s1=s2<<2`
 - Store in \$s1 the value from \$s2 shifted 2 bits to the left (they fall off end), inserting 0's on right; << in C.
- Before:
`0000 0002hex`
`0000 0000 0000 0000 0000 0000 0000 0010two`
- After:
`0000 0008hex`
`0000 0000 0000 0000 0000 0000 0000 1000two`
- What arithmetic effect does shift left have?
- Shift Right: `sr1` is opposite shift; >>

Arithmetic Shifting

- Shift right arithmetic moves n bits to the right (insert high order sign bit into empty bits)
- For example, if register `$s0` contained
 - `1111 1111 1111 1111 1111 1111 1110 0111`_{two} = `-25`_{ten}
- If executed `sra $s0, $s0, 4`, result is:
 - `1111 1111 1111 1111 1111 1111 1111 1110`_{two} = `-2`_{ten}
- Unfortunately, this is NOT same as dividing by 2^n
 - Fails for odd negative numbers
 - C arithmetic semantics is that division should round towards 0

And In Conclusion ...

- Computer words and vocabulary are called instructions and instruction set respectively
- MIPS is example RISC instruction set in this class
- Rigid format: 1 operation, 2 source operands, 1 destination
 - `add, sub, mul, div, and, or, sll, srl, sra`
 - `lw, sw, lb, sb` to move data to/from registers from/to memory.
 - On monday we will get control flow:
 - `beq, bne, j, slt, slti`
- Simple mappings from arithmetic expressions, array access, if-then-else in C to MIPS instructions