

# MIPS Functions and Instruction Formats

# The Contract: The MIPS Calling Convention

- You write functions, your compiler writes functions, other compilers write functions...
  - And all your functions call other functions which call other functions
- We want them all to play nicely together
- Thus the MIPS Calling Convention
  - How do you pass arguments?
  - What registers and state ***are not changed*** between when you call the function and when it returns?

# Why We Need it?

- ```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y;  
}
```
- What happens when a function calls another function?
  - Would clobber values in \$a0 to \$a3 and \$ra
- What is the solution?
  - Need to ***save values on the stack***

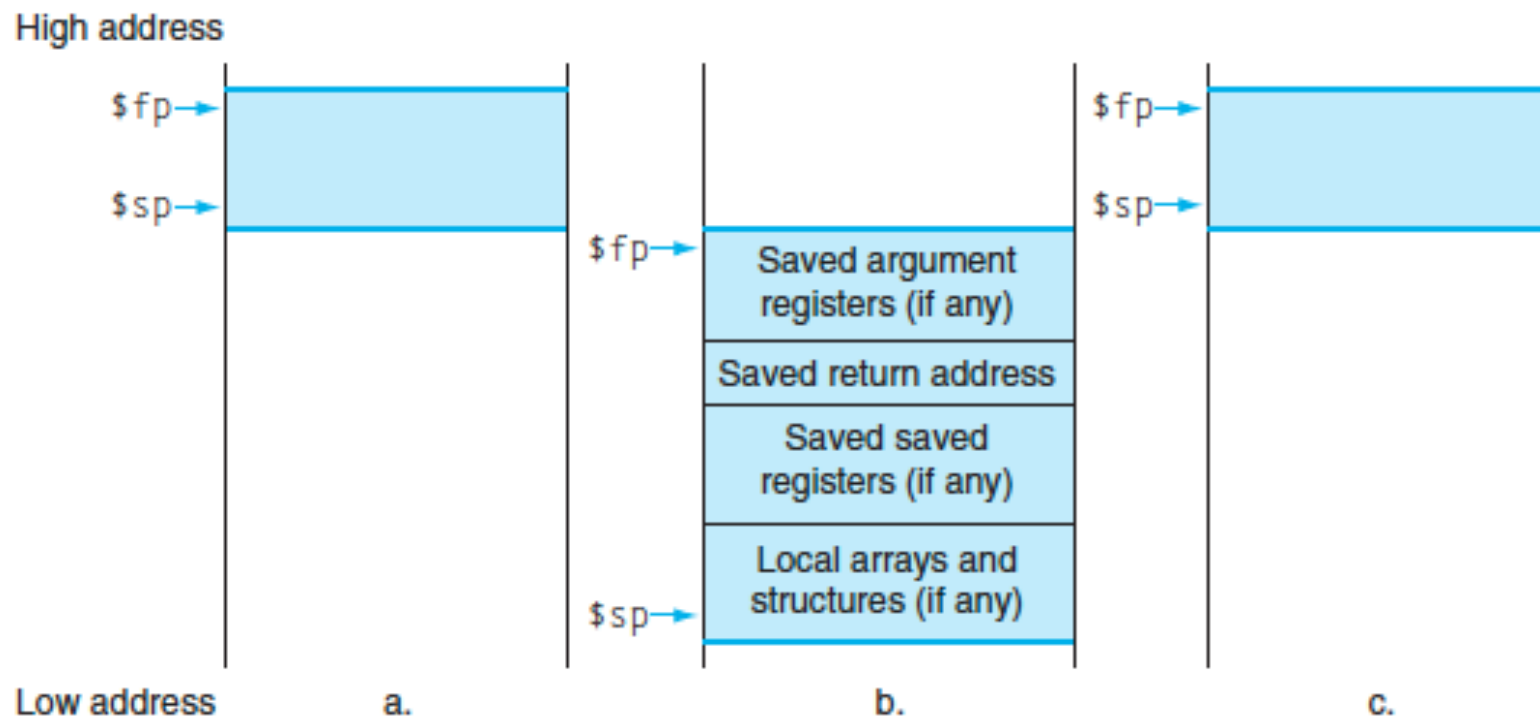
# Optimized Function Convention

- To reduce expensive loads and stores from spilling and restoring registers, MIPS divides registers into two categories
- Preserved across function call (Callee-saved)
  - Calling function can rely on values being unchanged when the called function returns
  - `$sp`, `$gp`, `$fp`, “saved registers” `$s0-$s7`
- Not preserved across function call (Caller-saved)
  - Caller cannot rely on values being unchanged, so the caller must save them if needed across a function call
  - Return value registers `$v0`, `$v1`, Argument registers `$a0-$a3`, `$t0-$t9`, `$ra`

# Allocating Space on Stack

- C has two storage classes: automatic and static
  - **Automatic** variables are local to function and discarded when function exits
  - Static variables exist across exits from and entries to procedures
- Use the stack for automatic (local) variables that don't fit in registers
- Use the stack for all callee-saved registers used in the function
  - And then restore those values upon function exit
- Procedure frame or activation record: segment of stack with saved registers and local variables
- Some MIPS compilers use a frame pointer (**\$fp**) to point to first word of frame
  - But in general we don't since, at any given point in the code, there is a fixed difference between the frame pointer and the stack pointer (**\$sp**)

# Stack Before, During, After Call



# Using the Stack (1/2)

- So we have a register `$sp` which always points to the last used space in the stack.
- To use the stack, we decrement this pointer by the amount of space we need and then fill it with info
  - And then when done, we restore anything that needs restoring before exit
- So, how do we compile this?
  - ```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y;  
}
```

# Using the Stack (2/2)

```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y; }
```

- Hand-compile

sumSquare:

```
“push”  addi $sp,$sp,-8 # space on stack  
        sw $ra, 4($sp)  # save ret addr  
        sw $a1, 0($sp)  # save y  
        add $a1,$a0,$zero # mult(x,x)  
        jal mult        # call mult  
        lw $a1, 0($sp)  # restore y  
        add $v0,$v0,$a1 # mult()+y  
        lw $ra, 4($sp)  # get ret addr  
“pop”   addi $sp,$sp,8  # restore stack  
        jr $ra  
mult:   ...
```



# Basic Structure of a Function


## Prologue

```
entry_label:  
addi $sp,$sp, -framesize  
sw $ra, framesize-4($sp)  # save $ra  
save other regs if need be
```

...  
**Body** (call other functions...)

## Epilogue

```
restore other regs if need be  
lw $ra, framesize-4($sp)  # restore $ra  
addi $sp,$sp, framesize  
jr $ra
```



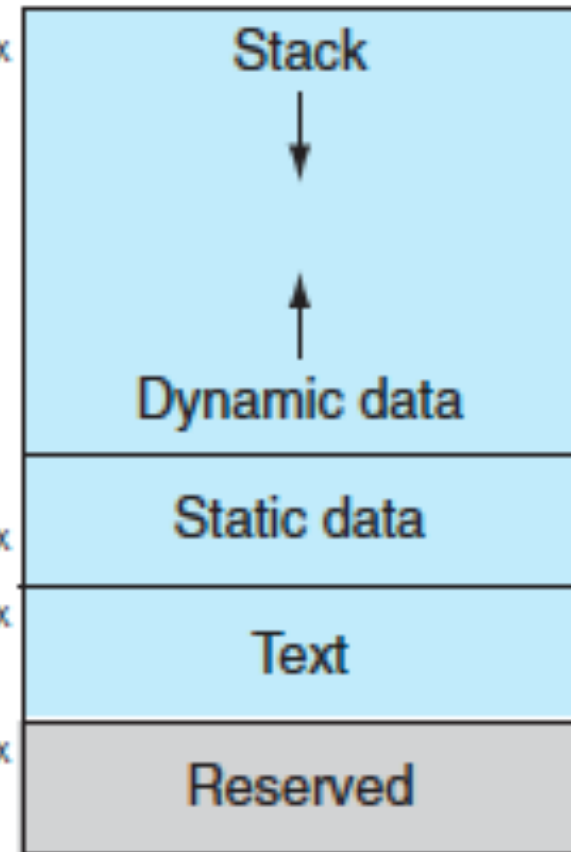
# MIPS Default Memory Allocation

- Text segment contains the code
- PC points to the start of it initially
- Stack grows down
  - `$sp` points to the lowest element
- `$gp` points to the start of static data

`$sp` → 7fff fffc<sub>hex</sub>

`$gp` → 1000 8000<sub>hex</sub>  
1000 0000<sub>hex</sub>

`pc` → 0040 0000<sub>hex</sub>  
0



# Register Allocation and Numbering

Name	Register number	Usage	Preserved on call?
\$zero	0	The constant value 0	n.a.
\$v0-\$v1	2-3	Values for results and expression evaluation	no
\$a0-\$a3	4-7	Arguments	no
\$t0-\$t7	8-15	Temporaries	no
\$s0-\$s7	16-23	Saved	yes
\$t8-\$t9	24-25	More temporaries	no
\$gp	28	Global pointer	yes
\$sp	29	Stack pointer	yes
\$fp	30	Frame pointer	yes
\$ra	31	Return address	yes

- \$1 reserved for assembler (\$at), \$26 and \$27 for the interrupt handler (\$k0 and \$k1)

# Recursive Function Factorial

```
int fact (int n)
{
    if (n < 1) return (1);
    else return (n * fact(n-1));
}
```

# Recursive Function Factorial

Fact:

```
# adjust stack for 2 items
addi $sp,$sp,-8
# save return address
sw $ra, 4($sp)
# save argument n
sw $a0, 0($sp)
# test for n < 1
slti $t0,$a0,1
# if n >= 1, go to L1
beq $t0,$zero,L1
# Then part (n==1) return 1
addi $v0,$zero,1
# pop 2 items off stack
addi $sp,$sp,8
# return to caller
```

jr \$ra

L1:

```
# Else part (n >= 1)
# arg. gets (n - 1)
addi $a0,$a0,-1
# call fact with (n - 1)
jal Fact
# return from jal: restore n
lw $a0, 0($sp)
# restore return address
lw $ra, 4($sp)
# adjust sp to pop 2 items
addi $sp, $sp,8
# return n * fact (n - 1)
mul $v0,$a0,$v0 mul is a pseudo instruction
# return to the caller
jr $ra
```

# The “Stack Overflow” Attack

- Recall that C allocates variables on the stack
  - And many c functions don't actually check that they are writing into valid memory...
- So what happens if you allocate an array on the stack...
  - And then call something that writes beyond the stack...

```
void foo() {  
    char bar[32] ;  
    ...  
    gets (bar) ;  
    ...  
}
```

# The Stack Overflow Continued...

```
foo:addi $sp $sp -36
    # allocate space for
    # $ra and bar
    sw $ra $sp(32)
...
    # bar == $sp + 0
    addi $a0 $sp $0
    jal gets
...
    lw $ra $sp(32)
    add $sp $sp 36
    jr $ra
```

...
...
\$ra
bar[28:31]
bar[24:27]
bar[20:23]
bar[16:19]
bar[12:15]
bar[8:11]
bar[4:7]
bar[0:3]
...

- Now what if a “bad dude” chooses the input into **gets**?
- Well, they can overwrite \$ra and also add their own code past that point...
- And make \$ra point to their own code...
- Voila, they’ve taken control!

# Oh, and jalr...

- We have **j**
  - "Jump to fixed address"
- **jr**
  - "Jump to address specified in register"
- **jal**
  - "Jump to this location and store PC+4 in \$ra"
- But we need one more: Jump and Link Register, **jalr**
  - "Jump to this register location and store PC+4 in \$ra"
- This is how we implement the pointers-to-functions ninjutsu!
  - `char (*f)(char *, char *) = &foo`

....

```
(*f)("arg1", "arg2") -> jalr $whateverFisStoredAt
```



# Clickers/Peer Instruction

- Which Statement is True?
  - a: `$sp` points to the lowest address currently in use on the stack
  - b: `$ra` stores PC+4 saved by `jal` so that a program knows where to return to when a function exits
  - c: `$t0-$t9` are callee saved registers
  - d: The classic stack overflow attack overwrites the saved return address on the stack with a new location
  - e: Nick likes trolling students on clicker questions


# EDSAC (Cambridge, 1949)

## First General Stored-Program Computer

Computer Science 61C Spring 2017

Friedland and Weaver

- Programs held as numbers in memory
- 35-bit binary 2's complement words



# Consequence #1: Everything Addressed

- Since all instructions and data are stored in memory, everything has a memory address: instructions, data words
  - both branches and jumps use these
- C pointers are just memory addresses: they can point to anything in memory
  - Unconstrained use of addresses can lead to nasty bugs; up to you in C; limited in Java by language design
- One register keeps address of instruction being executed: “Program Counter” (PC)
  - Basically a pointer to memory: Intel calls it Instruction Pointer (a better name)

# Consequence #2: Binary Compatibility

- Programs are distributed in binary form
  - Programs bound to specific instruction set
    - The "ABI": Application Binary Interface is a function of **both** the instruction set and the underlying operating system
  - Different binaries for Macintoshes and PCs
  - Different binaries for Linux i86 and Linux ARM
- New machines want to run old programs (“binaries”) as well as programs compiled to new instructions
  - Leads to “backward-compatible” instruction set evolving over time
- Selection of Intel 8086 in 1981 for 1st IBM PC is major reason latest PCs still use 80x86 instruction set;  
the hardware can still run programs from a 1981 PC today

# Instructions as Numbers (1/2)

- Currently all data we work with is in words (32-bit chunks):
  - Each register is a word.
  - `lw` and `sw` both access memory one word at a time.
- So how do we represent instructions?
  - Remember: Computer only understands 1s and 0s, so “`add $t0, $0, $0`” is meaningless.
  - MIPS/RISC seeks simplicity: since data is in words, make instructions be fixed-size 32-bit words also

# Instructions as Numbers (2/2)

- One word is 32 bits, so divide instruction word into “fields”.
- Each field tells processor something about instruction.
- We could define different fields for each instruction, but MIPS seeks simplicity, so define 3 basic types of instruction formats:
  - R-format
  - I-format
  - J-format

# Instruction Formats

- I-format: used for instructions with immediates, **lw** and **sw** (since offset counts as an immediate), and branches (**beq** and **bne**) since branches are "relative" to the PC
  - (but not the shift instructions)
- J-format: used for **j** and **jal**
- R-format: used for all other instructions
- It will soon become clear why the instructions have been partitioned in this way

# R-Format Instructions (1/5)

- Define “fields” of the following number of bits each:  $6 + 5 + 5 + 5 + 5 + 6 = 32$

6	5	5	5	5	6
---	---	---	---	---	---

- For simplicity, each field has a name:

<b>opcode</b>	<b>rs</b>	<b>rt</b>	<b>rd</b>	<b>shamt</b>	<b>funct</b>
---------------	-----------	-----------	-----------	--------------	--------------

- Important: On these slides and in book, each field is viewed as a 5- or 6-bit unsigned integer, not as part of a 32-bit integer
- Consequence: 5-bit fields can represent any number 0-31, while 6-bit fields can represent any number 0-63



# R-Format Instructions (2/5)

- What do these field integer values tell us?
  - opcode: partially specifies what instruction it is
    - Note: This number is equal to 0 for all R-Format instructions
  - funct: combined with opcode, this number exactly specifies the instruction
- Question: Why aren't **opcode** and **funct** a single 12-bit field?
  - We'll answer this later

# R-Format Instructions (3/5)

- More fields:
  - **rs** (Source Register): usually used to specify register containing first operand
  - **rt** (Target Register): usually used to specify register containing second operand (note that name is misleading)
  - **rd** (Destination Register): usually used to specify register which will receive result of computation

# R-Format Instructions (4/5)

- Notes about register fields:
  - Each register field is exactly 5 bits, which means that it can specify any unsigned integer in the range 0-31. Each of these fields specifies one of the 32 registers by number.
  - The word “usually” was used because there are exceptions that we’ll see later

# R-Format Instructions (5/5)

- Final field:
  - shamt: This field contains the amount a shift instruction will shift by. Shifting a 32-bit word by more than 31 is useless, so this field is only 5 bits (so it can represent the numbers 0-31)
  - This field is set to 0 in all but the shift instructions
- For a detailed description of field usage for each instruction, see green insert in COD
  - (We will provide a copy of the "green sheet" on all exams)

# R-Format Example (1/2)

- MIPS Instruction:
  - **add**     **\$8,\$9,\$10**
  - opcode = 0 (look up in table in book)
  - funct = 32 (look up in table in book)
  - rd = 8 (destination)
  - rs = 9 (first operand)
  - rt = 10 (second operand)
  - shamt = 0 (not a shift)

## R-Format Example (2/2)

opcode	rs	rt	rd	shamt	funct
--------	----	----	----	-------	-------

Computer Science 61C Spring 2017

Friedland and Weaver

- MIPS Instruction:

- add \$8,\$9,\$10
- Decimal number per field representation:

0	9	10	8	0	32
---	---	----	---	---	----

- Binary number per field representation:

000000	01001	01010	01000	00000	100000
--------	-------	-------	-------	-------	--------

- hex representation:

- 0x012A 4020

- Called a Machine Language Instruction

hex